

## **Low-Level Design Document (LLD)**

**Project Name:** Data Streaming (Kafka – NiFi – HDFS -Airflow– Hive Data Pipeline)

**Prepared By:** Marwa Mansour

**Date:** 11-Feb-2026

---

## **Code Details:**

### **Producer Application**

- Class Name: **KafkaAvroProducerApp**
- Responsibility:
  - Generates 1000 UserEvent records
  - Serializes records using Avro
  - Registers schema in Confluent Schema Registry
  - schema.registry.url=http://schema-registry:8081
  - Publishes records to Kafka topic: user-events

### **Consumer Application**

- Class Name: **KafkaAvroConsumerApp**
- Responsibility:
  - Consumes messages from Kafka topic user-events
  - Deserializes using KafkaAvroDeserializer
  - Prints GenericRecord output for validation
- **Topic Configuration**
  - • Topic Name: user-events
  - • Offset Reset: earliest
  - • Message Count Tested: 1000 records

## **Schema Registration:**

### **Avro Schema**

```
{  
  "type": "record",  
  "name": "UserEvent",  
  "namespace": "com.example.avro",  
  "fields": [  
    { "name": "user_id", "type": "int" },  
    { "name": "name", "type": "string" },  
    { "name": "email", "type": "string" },  
    { "name": "event_type", "type": "string" },  
    { "name": "event_time", "type": "long" }  
  ]  
}
```

## KafkaAvroProducerApp Code To Generate 1000 Record:

```
package com.example;
import com.example.avro.UserEvent;
import org.apache.kafka.clients.producer.*;
import io.confluent.kafka.serializers.KafkaAvroSerializer;
import java.util.Properties;
public class KafkaAvroProducerApp {

    public static void main(String[] args) {

        String bootstrapServers = "localhost:9092";
        String schemaRegistryUrl = "http://localhost:8081";
        String topic = "user-events";

        Properties props = new Properties();
        props.put(ProducerConfig.BOOTSTRAP_SERVERS_CONFIG, bootstrapServers);
        props.put(ProducerConfig.KEY_SERIALIZER_CLASS_CONFIG, KafkaAvroSerializer.class.getName());
        props.put(ProducerConfig.VALUE_SERIALIZER_CLASS_CONFIG, KafkaAvroSerializer.class.getName());
        props.put("schema.registry.url", schemaRegistryUrl);

        Producer<String, UserEvent> producer = new KafkaProducer<>(props);

        for (int i = 1; i <= 1000; i++) {

            UserEvent userEvent = UserEvent.newBuilder()
                .setId(i)
                .setName("User" + i)
                .setEmail("user" + i + "@example.com")
                .setEventTime(System.currentTimeMillis())
                .build();

            ProducerRecord<String, UserEvent> record =
                new ProducerRecord<>(topic, String.valueOf(userEvent.getId()), userEvent);

            producer.send(record, (metadata, exception) ->{
                if (exception != null) {
                    exception.printStackTrace();
                } else {
                    System.out.println("Sent record " + userEvent.getId() +
                        " to partition " + metadata.partition() +
                        " with offset " + metadata.offset());
                }
            });
        }

        producer.flush();
        producer.close();

        System.out.println("All 1000 messages sent successfully!");
    }
}
```

## **KafkaAvroConsumerApp To consumer data record from producer:**

```
package com.example;
import org.apache.kafka.clients.consumer.*;
import io.confluent.kafka.serializers.KafkaAvroDeserializer;
import org.apache.avro.generic.GenericRecord;
import java.time.Duration;
import java.util.Arrays;
import java.util.Properties;

public class KafkaAvroConsumerApp {
    public static void main(String[] args) {
        Properties props = new Properties();
        props.put("bootstrap.servers", "localhost:9092");
        props.put("group.id", "test-group");
        props.put("key.deserializer", KafkaAvroDeserializer.class.getName());
        props.put("value.deserializer", KafkaAvroDeserializer.class.getName());
        props.put("schema.registry.url", "http://localhost:8081");
        props.put("specific.avro.reader", "false");

        KafkaConsumer<String, GenericRecord> consumer = new KafkaConsumer<>(props);
        consumer.subscribe(Arrays.asList("user-events"));

        while (true) {
            ConsumerRecords<String, GenericRecord> records = consumer.poll(Duration.ofMillis(1000));
            for (ConsumerRecord<String, GenericRecord> record : records) {
                System.out.println(record.value());
            }
        }
    }
}
```

## **NiFi Flow Detailed Design (Production Level)**

### **Flow Name:**

kafka\_nifi\_hdfs\_hive

### **Purpose**

This NiFi flow is designed to consume real-time streaming data from Kafka, validate and enrich the records, persist them in HDFS in a structured format, and provide full observability through structured logging and failure handling mechanisms.

The flow follows a controlled success/failure routing strategy to ensure reliability, data integrity, and production-grade monitoring.

---

## **ConsumeKafkaRecord\_2\_6**

### **Component Role**

Acts as the ingestion entry point of the streaming pipeline.

- **Functional Responsibilities**
- Subscribes to Kafka topic: user-events
- Deserializes Avro messages using configured Record Reader
- Converts Kafka records into NiFi FlowFiles
- Commits offsets only after successful processing (consumer group controlled)

<b>Relationship</b>	<b>Action</b>
success	Routes valid records to UpdateAttribute
parse.failure	Routes malformed or schema-invalid records to logging processor

### **Production Considerations**

- Consumer group configured for controlled offset tracking
- Back pressure enabled to prevent memory overflow
- Supports horizontal scaling when deployed in NiFi cluster

---

## **2. Parse Failure Handling (LogAttribute – Failure Path)**

### **Purpose**

Captures malformed or schema-incompatible records.

### **Observability**

This ensures early detection of schema incompatibility or producer-side data issues.

---

## **UpdateAttribute**

### **Component Role**

Prepares the FlowFile for structured storage.

### **Functional Responsibilities**

- Generates dynamic filename if required
- Adds metadata attributes for lineage tracking

### **Design Rationale**

- Enables dynamic HDFS directory structure.

---

## **PutHDFS**

### **Component Role**

Handles distributed storage into HDFS.

### **Functional Responsibilities**

- Writes records in Parquet format
- Stores data under directory structure
- Ensures compatibility with Hive external table

### **Target Directory Structure**

/data/user\_events/ user\_events\_<timestamp>.parquet

### **Relationships**

<b>Relationship</b>	<b>Action</b>
success	Routed to Success LogAttribute
failure	Routed to Failure LogAttribute

### **Production Considerations**

- HDFS replication factor ensures durability
  - Supports large-scale distributed storage
  - Fault-tolerant write mechanism
-

## End-to-End Execution Lifecycle

1. Kafka message ingestion
  2. Schema-based deserialization
  3. Validation routing (success vs parse failure)
  4. Attribute enrichment
  5. Parquet file creation
  6. HDFS persistence
  7. Success/Failure logging
- 

## Reliability & Enterprise Controls

### **✓ Controlled Offset Management**

Offsets committed only after successful processing.

### **✓ Schema Governance**

Strict Avro validation prevents malformed ingestion.

### **✓ Observability**

Logging implemented at ingestion, transformation, and storage stages.

### **✓ Failure Isolation**

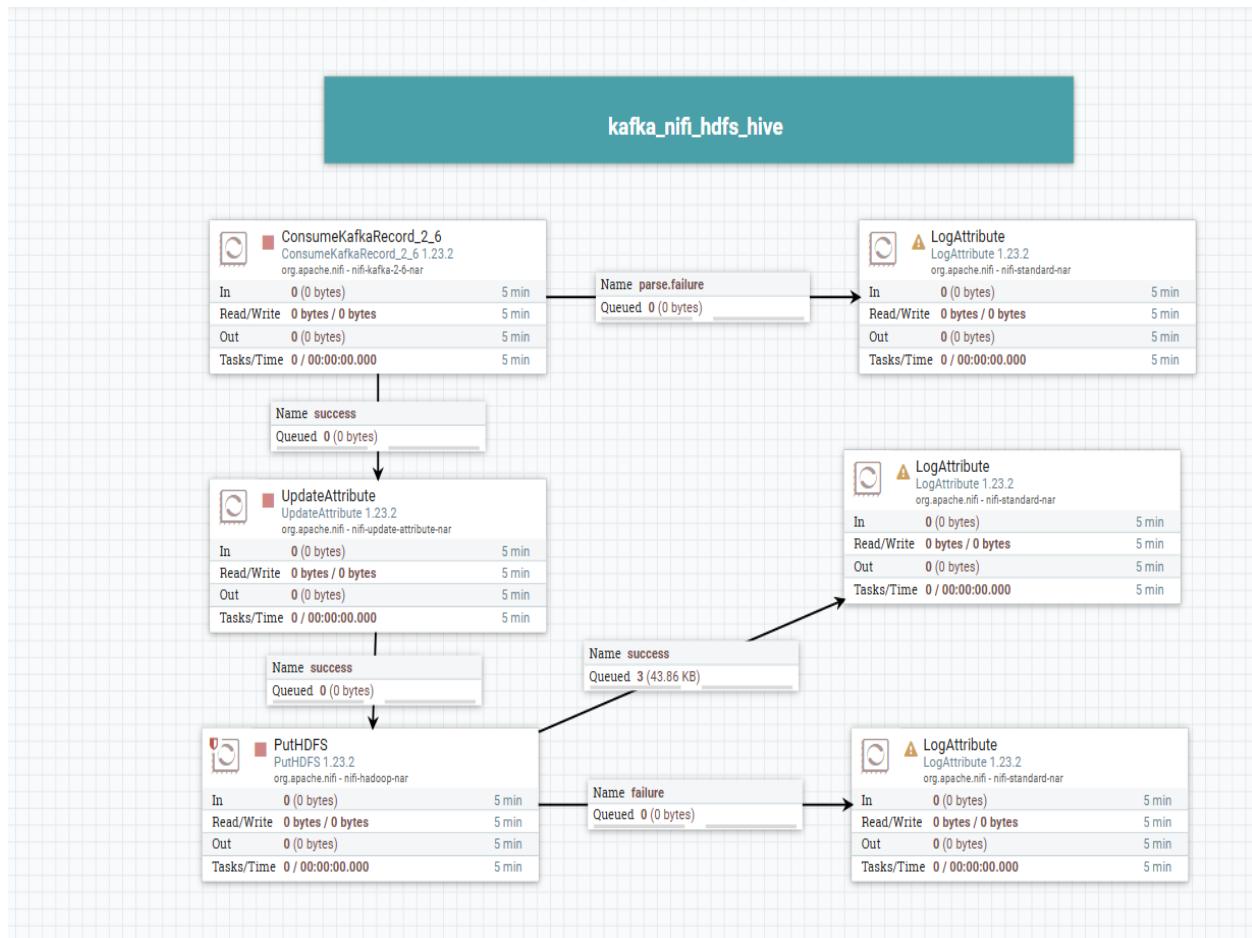
Corrupted records are isolated and do not impact healthy data flow.

---

## Production Readiness Summary

This NiFi flow design ensures:

- High reliability
- Clear data lineage
- Strong schema enforcement
- Distributed scalable storage
- Enterprise-grade monitoring
- Fault-tolerant streaming ingestion



## Connection Details

- **Kafka:**
  - Broker: kafka:29092
  - Topic: user-events
  - Schema Registry: <http://schema-registry:8081>
- **NiFi:**
  - URL: http://localhost:8060
  - Processor: ConsumeKafkaRecord\_2\_6
  - Record Reader: AvroReader
  - Record Writer: ParquetRecordSetWriter
- **HDFS:**
  - Path: /data/user\_events
- **Hive:**
  - Database: user\_events\_db



## DDLs (Hive Tables)

### ○ External Table (Parquet)

```
CREATE EXTERNAL TABLE user_events_db.user_events_parq (
    id INT,
    name STRING,
    email STRING,
    event_time BIGINT
)
STORED AS PARQUET
LOCATION '/data/user_events';
```

### ○ Managed Table (Parquet)

```
CREATE TABLE user_events_prod.user_events_parquet_managed (
    id INT,
    name STRING,
    email STRING,
    event_time BIGINT
)
STORED AS PARQUET
TBLPROPERTIES ("parquet.compression"="SNAPPY");
```

### ○ Load data into managed table :

#### ○ Method 1:

```
LOAD DATA INPATH '/data/user_events' INTO TABLE user_events_managed;
```

#### ○ Methode 2 (loading data from external table into managed table):

```
INSERT INTO TABLE user_events_parquet_managed
```

```
SELECT * FROM user_events_db.user_events_parq;
```

## Data Loading & Archiving Logic

### 1 Overview

This section describes the detailed implementation of the data ingestion process from HDFS to Hive managed table using Apache Airflow.

The pipeline ensures:

- Only new files are loaded into the Hive managed table.
- Previously processed files are archived.
- Data is stored in Parquet format with Snappy compression.
- Duplicate loading is prevented.

---

### 2 HDFS Directory Structure

/data/user\_events/ → Incoming files

/data/user\_events/archive/ → Archived processed files

---

### 3 Hive Configuration

- Database: user\_events\_prod
- Managed Table: user\_events\_parquet\_managed
- Storage Format: Parquet
- Compression: Snappy

Table type: Managed Table(id , name , email, event\_time)

Location: Managed by Hive (default warehouse path)

---

### 4 Processing Logic

#### Step 1: Detect New Files

Airflow checks the HDFS incoming directory.

#### Step 2: Load to Hive Managed Table

New files are inserted into:

user\_events\_prod.user\_events\_parquet\_managed

Using:

- Parquet format
- Snappy compression

### **Step 3: Archive Processed Files**

After successful load:

- Files are moved from:
- /data/user\_events/

To:

/data/user\_events/archive/

This ensures:

- Idempotent processing
  - No duplicate ingestion
  - Clear separation between new and processed data
- 

### **5 DAG Implementation**

DAG Name: user\_events\_kafka\_nifi\_hadoop\_hive

Main Task:

- load\_and\_archive

Execution Flow:

1. Load new data into Hive
  2. Move processed files to archive
  3. Log execution status
- 

### **6 Error Handling**

- If Hive load fails → files are NOT moved to archive.
- If archive step fails → DAG retries based on retry configuration.
- Logs are stored in Airflow logs directory.

 **Airflow Dag Code:**

```
from airflow import DAG
from airflow.operators.bash import BashOperator
from datetime import datetime, timedelta
# Paths and Hive table names
HDFS_PATH = "/data/user_events"
ARCHIVE_PATH = "/data/user_events/archive"
HIVE_DB = "user_events_prod"
HIVE_TABLE = "user_events_parquet_managed"

# Default arguments for the DAG
default_args = {
    'owner': 'airflow',
    'start_date': datetime(2026, 2, 16),
    'retries': 1,
    'retry_delay': timedelta(minutes=2),
}

# Define the DAG
with DAG(
    dag_id="user_events_kafka_nifi_hadoop_hive",
    default_args=default_args,
    schedule_interval="*/5 * * * *",
    catchup=False,
    max_active_runs=1,
    tags=['hive', 'hdfs', 'archive'],
) as dag:

    # Task: Load new files into managed Hive table and archive them
    load_and_archive = BashOperator(
        task_id="load_and_archive",
        bash_command=f"""
            set -e
            echo "Creating archive directory if not exists..."
            hdfs dfs -mkdir -p {ARCHIVE_PATH}
        """
    )

    # Count new files
    file_count=$(hdfs dfs -ls {HDFS_PATH} 2>/dev/null | grep -v '/archive' | grep 'user-events-*\\.parquet$' | wc -l)

    if [ "$file_count" -eq 0 ]; then
```

```

echo "No new files to process"
exit 0
fi

echo "Found $file_count new file(s)"

temp_ext_table="ext_user_events_temp_$(date +%s)"
echo "Creating temporary external Hive table: $temp_ext_table"

hive -e "
USE {HIVE_DB};
SET hive.stats.autogather=false;
SET hive.compute.query.using.stats=false;

CREATE EXTERNAL TABLE $temp_ext_table (
    id INT,
    name STRING,
    email STRING,
    event_time BIGINT
)
STORED AS PARQUET
LOCATION '{HDFS_PATH}';

INSERT INTO TABLE {HIVE_TABLE}
SELECT t.id, t.name, t.email,t.event_time
FROM $temp_ext_table t
LEFT JOIN {HIVE_TABLE} h ON t.id = h.id
WHERE h.id IS NULL;

DROP TABLE $temp_ext_table;
"
echo "? Data loaded to managed table successfully"
# Archive processed files
for file in $(hdfs dfs -ls {HDFS_PATH} | awk '{{print $8}}' | grep 'user-events-*\\.parquet$'); do
    filename=$(basename "$file")
    hdfs dfs -mv "$file" "{ARCHIVE_PATH}/$filename"
    echo "? Archived: $filename"
done
echo "? All files archived successfully"
""",

```